Free Questions for Databricks-Certified-Data-Engineer-Associate

Shared by Randolph on 05-08-2025

For More Free Questions and Preparation Resources

Check the Links on Last Page

Question 1

Question Type: MultipleChoice

A data engineer needs to use a Delta table as part of a data pipeline, but they do not know if they have the appropriate permissions.

In which location can the data engineer review their permissions on the table?

Options:

A- Jobs
B- Dashboards
C- Catalog Explorer
D- Repos

Answer:

C

Question 2

Question Type: MultipleChoice

Which type of workloads are compatible with Auto Loader?

Options:

A- Streaming workloads
B- Machine learning workloads
C- Serverless workloads
D- Batch workloads

Answer:

A

Question 3

Question Type: MultipleChoice

Which SQL keyword can be used to convert a table from a long format to a wide format?

Options:

A- TRANSFORM
B- PIVOT
C- SUM
D- CONVERT

Answer:

B

Question 4

Question Type: MultipleChoice

Which two components function in the DB platform architecture's control plane? (Choose two.)

Options:

A- Virtual Machines
B- Compute Orchestration
C- Serverless Compute
D- Compute
E- Unity Catalog

Answer:

B, E

Question 5

Question Type: MultipleChoice

Identify the impact of ON VIOLATION DROP ROW and ON VIOLATION FAIL UPDATE for a constraint violation.

A data engineer has created an ETL pipeline using Delta Live table to manage their company travel reimbursement detail, they want to ensure that the if the location details has not been

provided by the employee, the pipeline needs to be terminated.

How can the scenario be implemented?

Options:

A- CONSTRAINT valid_location EXPECT (location = NULL)
B- CONSTRAINT valid_location EXPECT (location != NULL) ON VIOLATION FAIL UPDATE
C- CONSTRAINT valid_location EXPECT (location != NULL) ON DROP ROW
D- CONSTRAINT valid_location EXPECT (location != NULL) ON VIOLATION FAIL

Answer:

B

Question 6

Question Type: MultipleChoice

Which method should a Data Engineer apply to ensure Workflows are being triggered on schedule?

Options:

A- Scheduled Workflows require an always-running cluster, which is more expensive but reduces processing latency.
B- Scheduled Workflows process data as it arrives at configured sources.
C- Scheduled Workflows can reduce resource consumption and expense since the cluster runs only long enough to execute the pipeline.
D- Scheduled Workflows run continuously until manually stopped.

Answer:

C

Question 7

Question Type: MultipleChoice

Identify a scenario to use an external table.

A Data Engineer needs to create a parquet bronze table and wants to ensure that it gets stored in a specific path in an external location.

Which table can be created in this scenario?

Options:

A- An external table where the location is pointing to specific path in external location.
B- An external table where the schema has managed location pointing to specific path in external location.
C- A managed table where the catalog has managed location pointing to specific path in external location.
D- A managed table where the location is pointing to specific path in external location.

Answer:

A

Question 8

Question Type: MultipleChoice

Identify how the count_if function and the count where x is null can be used

Consider a table random_values with below data.

What would be the output of below query?

select count_if(col > 1) as count_

a. count(*) as count_b.count(col1) as count_c from random_values col1

0

1

2

NULL -

2

3

Options:

A- 3 6 5
B- 4 6 5
C- 3 6 6
D- 4 6 6

Answer:

A

Question 9

Question Type: MultipleChoice

A data engineer has configured a Structured Streaming job to read from a table, manipulate the data, and then perform a streaming write into a new table.

```
(spark.table("sales")
    .withColumn("avg_price", col("sales") / col("units"))
    .writeStream
    .option("checkpointLocation", checkpointPath)
    .outputMode("complete")
    ._____
    .table("new_sales")
)
```

The code block used by the data engineer is below:

Which line of code should the data engineer use to fill in the blank if the data engineer only wants the query to execute a micro-batch to process data every 5 seconds?

Options:

A- trigger('5 seconds')
B- trigger(continuous='5 seconds')
C- trigger(once='5 seconds')
D- trigger(processingTime='5 seconds')

Answer:

D

Question 10

Question Type: MultipleChoice

The Delta transaction log for the 'students' tables is shown using the 'DESCRIBE HISTORY students' command. A Data Engineer needs to query the table as it existed before the UPDATE operation listed in the log.

	version	timestamp	operation
1	8	2024–04–22T14:33:31.000	OPTIMIZE
2	7	2024–04–22T14:33:16.000	MERGE
3	6	2024–04–22T14:33:06.000	DELETE
4	5	2024–04–22T14:32:58.000	UPDATE
5	4	2024–04–22T14:32:47.000	WRITE
6	3	2024–04–22T14:32:44.000	WRITE
7	2	2024–04–22T14:32:23.000	WRITE
8	1	2024–04–22T14:32:20.000	WRITE
9	0	2024–04–22T14:31:49.000	CREATE TABLE

Which command should the Data Engineer use to achieve this? (Choose two.)

Options:

A- SELECT * FROM students@v4

B- SELECT * FROM students TIMESTAMP AS OF '2024-04-22T 14:32:47.000+00:00'

C- SELECT * FROM students FROM HISTORY VERSION AS OF 3

D- SELECT * FROM students VERSION AS OF 5

E- SELECT * FROM students TIMESTAMP AS OF '2024-04-22T 14:32:58.000+00:00'

Answer:

A, B

To Get Premium Files for Databricks-Certified-Data-Engineer-Associate Visit
https://www.p2pexams.com/products/databricks-certified-data-engineer-associate

For More Free Questions Visit
https://www.p2pexams.com/databricks/pdf/databricks-certified-data-engineer-associate

20% DISCOUNT